

Get started

- Overview
- Quickstart
- Models
- Pricing
- SDKs and CLI
- Latest: GPT-5.5
- Prompt guidance

Core concepts

- Text generation
- Code generation
- Images and vision
- Audio and speech
- Structured output
- Function calling
- Responses API
- Using tools

Agents SDK

- Overview
- Quickstart
- Agent definitions
- Models and providers
- Running agents
- Sandbox agents
- Orchestration
- Guardrails
- Results and state
- Integrations and observability
- Evaluate agent workflows
- Voice agents
- Agent Builder

Tools

- Web search
- MCP and Connectors
- Skills
- Shell
- Computer use
- File search and retrieval
- Tool search
- More tools

Run and scale

- Conversation state
- Background mode
- Streaming
- WebSocket mode
- Webhooks
- File inputs
- Context management
- Prompting
- Reasoning

Evaluation

- Getting started
- Working with evals
- Prompt optimizer

Evaluate agent workflows

Use traces, graders, datasets, and eval runs to improve agent quality.

Copy Page

The OpenAI Platform offers a suite of evaluation tools to help you ensure your agents perform consistently and accurately.

Use this page as the decision point for the evaluation surfaces that matter most for agent workflows.

Start with traces when you are still debugging behavior

Trace grading is the fastest way to identify workflow-level issues. A trace captures the end-to-end record of model calls, tool calls, guardrails, and handoffs for one run. Graders let you score those traces with structured criteria so you can find regressions and failure modes at scale.

Use trace grading when you want to answer questions like:

- Did the agent pick the right tool?
- Did a handoff happen when it should have?
- Did the workflow violate an instruction or safety policy?
- Did a prompt or routing change improve the end-to-end behavior?

Trace-grading workflow

- Open **Logs > Traces** in the dashboard.
- Inspect a representative workflow trace from Agent Builder or an SDK-based app with tracing enabled.
- Create a grader and run it against the selected traces.
- Use the results to refine prompts, tool surfaces, routing logic, or guardrails.

For code-first SDK workflows, start with [Integrations and observability](#) to get high-signal traces before you formalize graders.

Move to datasets and eval runs when you need repeatability

Once you know what “good” looks like, move from individual traces to repeatable datasets and eval runs. This is the right step when you want to benchmark changes, compare prompts, or run larger-scale evaluations over time.

If you need advanced features such as evaluation against external models, evaluation APIs, or larger-scale batch evaluation, use [Evals](#) alongside datasets.

Related evaluation surfaces

Getting started with evals: Datasets
Operate a flywheel of continuous improvement using evaluations.

Working with evals
Evaluate against external models, interact with evals via API, and more.

Prompt optimizer
Use your dataset to automatically improve your prompts.

Cookbook: Building resilient prompts with evals
Operate a flywheel of continuous improvement using evaluations.

Start with traces when you are still debugging behavior

Move to datasets and eval runs when you need repeatability

Related evaluation surfaces

Copy Page